GHOST: Generalizing Skills By Hierarchical Learning from Human Demonstrations

Sriram Krishna[1], Ben Eisner[1], Haotian Zhan[1], Ying Yuan[1], Haoyu Zhen[2],
Chuang Gan[2], Shubham Tulsiani[1], David Held[1]
[1]Robotics Institute, Carnegie Mellon University
[2]UMass Amherst

https://ghost-human-demo.github.io/

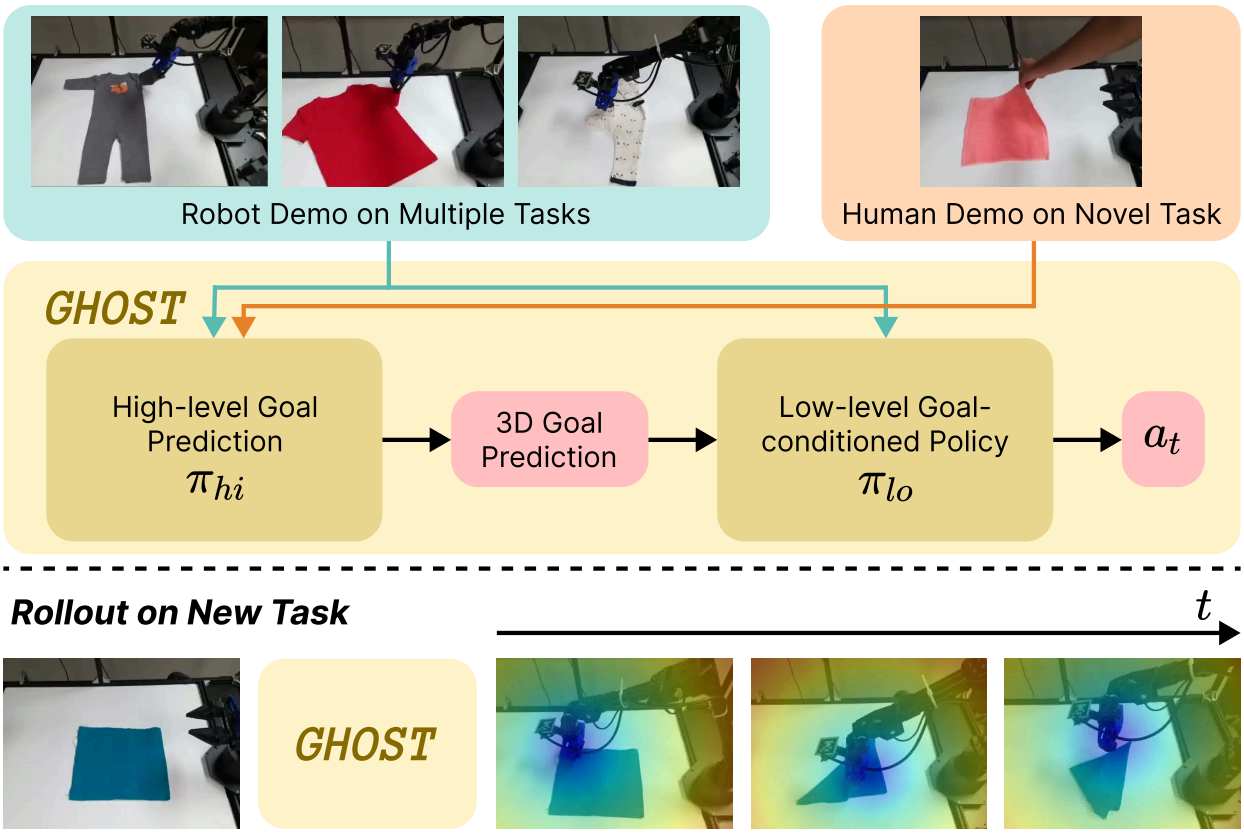

Fig. 1: GHOST learns robot skills from in-distribution teleoperation data and generalizes to novel tasks using human demonstrations. We train a hierarchical policy that decouples task-specific goal prediction (π_{hi}) from embodiment-specific action execution (π_{lo}). By training π_{hi} on both robot and human data while keeping π_{lo} purely on robot data, GHOST transfers learned manipulation skills to out-of-distribution tasks with only human video demonstrations.

***Abstract*—We present GHOST, a framework for generalizing robot skills to new tasks through human demonstrations. Robot teleoperation data is critical for training robust policies, but is expensive to scale, while human video is abundant, but lacks executable robot actions. We use a hierarchical policy framework that decouples a policy into two parts. The high-level policy maps visual observations to task "sub-goals". These sub-goals are intermediate objectives like "grasp the handle" defined in a task-relevant but embodiment-agnostic manner, and do not require action labels. The low-level policy is conditioned on these goals to execute embodiment-specific actions and is trained purely on robot data. This setup provides two key benefits. First, even when trained purely on teleoperation data, it significantly boosts performance. Second, it enables out-of-distribution generalization through human demonstrations - allowing us to learn precise embodiment-specific skills from robot data and then *generalize* these learned skills across novel objects and compositions of these skills.**

I. Introduction

The dominant paradigm today for teaching robot policies to manipulate their environment is imitation learning (IL). Given a sufficiently large number of demonstration trajectories, imitation learning algorithms can learn precise, complex behaviors [5, 34]. However, to generalize to novel objects and environments, they require massive amounts of data, typically

collected through *teleoperation*, where a human operator controls a robot to complete a given task [3, 11, 1]. This recipe, while simple, is expensive and labor-intensive to scale.

Human demonstrations are a promising way to collect large amounts of data much cheaper and faster, either through the use of web videos of humans interacting naturally with their environment [7, 8], or through task-specific demonstrations [12]. However, human data is *passive i.e.,* it does not provide any action labels with which to train imitation learning algorithms. Recent works use human demonstrations in many ways - one line of research uses vast amounts of web-scale data to learn general motion priors, and then finetune on a small amount of robot data (a few hundred trajectories) to "specialize" it to a specific embodiment [6, 2, 29]. Although these methods exhibit strong semantic generalization, they are often limited in terms of skill *dexterity*, as they cannot extract clean action labels from human web videos. Another line of work uses task-specific demonstrations collected by a human and *retargets* the human hand - extracting an action label from the human demo, which is then used to train a policy [12, 21, 9]. The most common instantiation of this paradigm involves using off-the-shelf hand pose estimators [20, 18] and mapping the detected hand keypoints to those of the robot end effector. However, using these "pseudo-actions" to train the policy depends on quality of retargeting, and poor retargeting can cause the policy to exhibit undesirable behavior - *an embodiment gap.*

We consider the setting where we have demonstrations on a set of objects and tasks (**in-distribution**) but our policy needs to be deployed on a novel task / object for which robot data is not available, but human demonstrations are (**out-of-distribution (OOD)**). Prior work looks at *enhancing* performance by supplementing small amounts of robot data with human demonstrations on the same task, while we consider out-of-distribution generalization. Assuming access to human demonstrations for such a deployment scenario, how do we adapt our policy to generalize its learned skills?

We approach this problem through the lens of a hierarchical policy framework. We decompose the policy into high-level goal prediction and low-level goal-conditioned policy [26]. We train a high-level goal prediction network that learns to predict the spatial pose of the end-effector at the end of a sub-goal, whereas the low-level policy is conditioned on the predicted goals to execute actions. We extend standard RGB-based imitation learning algorithms to condition on these 3D goals by projecting them to the image plane and representing them as *end-effector heatmaps*. **Our key insight is that this hierarchical factorization with spatial goals enables the high-level policy to learn *task-specific embodiment-agnostic* goals, whereas the low-level goal-conditioned policy learns *task-agnostic embodiment-specific* skills**. Critically, we can train the high-level policy on robot/human video as goals are embodiment-agnostic, whereas the low-level policy is trained purely on robot data to learn precise manipulation and goal-following behavior.

We find that this factorization yields two benefits: First, we show that even with only robot data, the hierarchical policy framework improves performance on in-distribution tasks. Second, through human demonstrations, we show that the framework improves OOD performance as the goal prediction transfers whereas action execution is unaffected by the human demonstrations.

Thus, we present **GHOST**, a hierarchical policy learning framework for training visuomotor policies that can generalize skills through human demonstrations. GHOST features: a) a 3D high-level goal prediction network that takes in RGB-D images to predict a distribution of sub-goal end-effector poses. b) a low-level policy conditioned on goals represented as heatmap projections of the predicted end-effector sub-goals.

Our contributions are as follows:

- A hierarchical framework for training policies on heterogeneous data sources that improves performance on in-distribution tasks and generalizes to out-of-distribution variations of the same tasks.
- A novel architecture for 3D goal prediction, aggregating RGB-D information to model distributions of sub-goals.
- A cross-embodiment spatial goal representation using *end-effector heatmaps* that enables effective goal-conditioning of image-based policies.

II. Related Work

A. Learning from Demonstrations

1) Imitation Learning: : Imitation Learning (IL) is a framework in which an agent is trained to perform sequential decision-making tasks by mimicking the behavior of an expert, from whom demonstrations are provided. While conceptually simple, it has gained widespread popularity in robotics due to the surprising complexity of behaviors that can be learned, given enough data. Recent work in this direction has been based on developing deep neural architectures that can learn dexterous and reactive policies from raw demonstration data, such as Zhao *et al.* [34], and Chi *et al.* [5] and Ze *et al.* [31]. When scaled up through large dataset collection endeavors [11, 17], much larger models can be trained and exhibit strong generalization behavior [4, 3]. However, these methods are fundamentally bottlenecked by the requirement of a large demonstration dataset, which can be unsafe, labor-intensive and expensive to collect at scale.

2) Hierarchical Imitation Learning: : Hierarchical Imitation Learning factorizes the IL framework into a high-level "planner" and a low-level "control" policy. Previous work has primarily focused on learning high- and low-level policies from robot data [15, 28, 14]. More recently, Wang *et al.* [26] demonstrated the effectiveness of a hierarchical policy to generalize to numerous articulated objects by representing sub-goals as end-effector poses. However, they do so only in the context of a single task (articulated object opening) and rely only on geometry as input to maximize sim2real transfer. We demonstrate a more general framework that can be applied to arbitrary tasks. MimicPlay [25] also use a hierarchical imitation learning framework where the high-level policy is

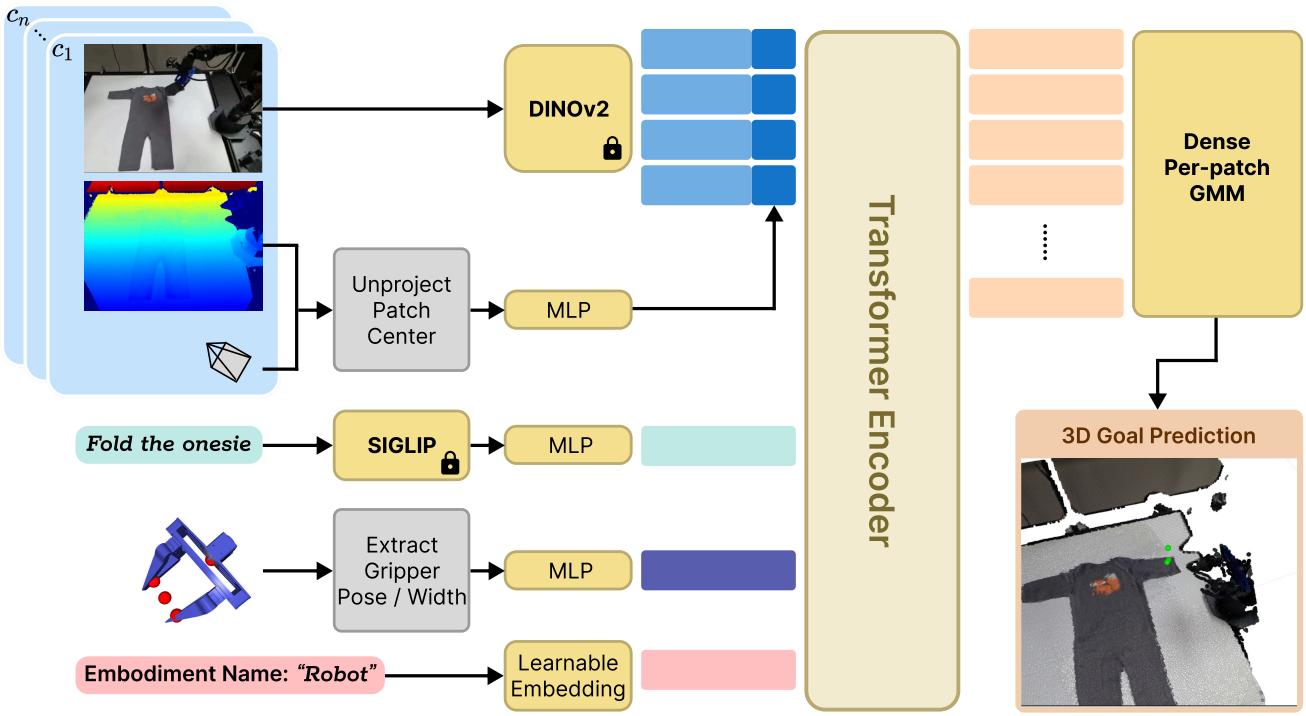

Fig. 2: **GHOST** High-level goal prediction architecture: RGB-D observations from multiple cameras are processed with a DINOv2 encoder, with patch tokens augmented by 3D coordinates. Additional context (gripper state, language embedding, embodiment name) is encoded via separate MLPs. An encoder-only transformer processes all tokens, with each patch predicting GMM parameters of a 3D goal distribution over end-effector keypoints.

learned from human play data to learn a latent plan over 3D end-effector trajectories. In contrast, we represent our goals as discrete end-effector goal configurations and use task-specific human demonstrations instead of play data.

B. Learning from Human Demonstrations

1) Learning policies: : With the rapid improvement in the accuracy of hand pose estimators [18, 20, 30], one direction of work focuses on tracking hand motions in human video demonstrations and explicitly representing them as robot actions in end-effector space, *i.e., retargeting* the human hand to the end-effector pose of the robot. Kareer *et al.* [10] train a unified human-robot policy by modifying ACT [34] to predict both joint positions (for robot data) and end-effector poses (for human and robot data). Ren *et al.* [21] and Haldar and Pinto [9] learn keypoint prediction models from human and/or robot demonstrations, and at inference, these are retargeted to an end-effector action based on rigid body constraints. While Haldar and Pinto [9] overcome the visual human-robot domain gap by operating purely in point-space, Lepert, Fang and Bohg *et al.* [12] overlay a virtual robot onto the human arm to minimize the visual gap when training a policy from retargeted human data. As these methods aim to directly learn a policy from human demonstrations, they use minimal (or zero) robot data, but rely on the quality of hand pose estimation, and that the collected human data consists of trajectories that are executable by the robot.

2) Learning goals: : On the other hand, another line of work aims to learn "goals" and "plans" from human video, which are then used to condition a policy (trained on robot data) to guide its behavior. Bharadhwaj *et al.* [2] and Wen *et al.* [27] learn explicit track prediction models that learn to predict future trajectories of pixels conditioned on a goal image/language instruction. These track predictions are then used to condition the robot control policies to execute a given task. Collins *et al.* [6] and Ye *et al.* [29] learn *latent* dynamics from large amounts of web-scale video, with the former learning latent tracks and the latter learning latent actions through a VQ-VAE [24]. These latent dynamics are then mapped to robot actions through a separate inverse dynamics model, trained with a small amount of robot demonstrations. Our work is closest to MimicPlay [25], which learns a latent distribution over hand trajectories collected through play data, with a key distinction being that they rely on a video prompt of the task being provided at test-time, whereas we only require a language goal to be provided.

III. Problem Statement and Assumptions

Task Definition. We consider manipulation tasks that can be decomposed into sequential phases (*sub-goals*). We denote a *skill* as a manipulation capability described through a language instruction (*e.g.,* grasping, placing, folding) that can be instantiated across different objects and contexts [4]. Each task consists of one or more sub-goals representing transitions in

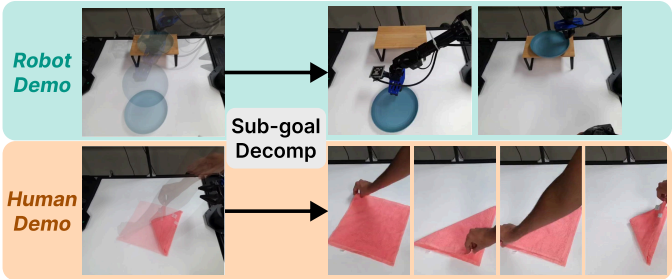

Fig. 3: Sub-goal decomposition for robot and human demonstrations. **Top**: Robot teleop demonstrations with sub-goals automatically extracted at gripper state transitions. **Bottom**: Human demonstrations with manually annotated sub-goals. Each sub-goal represents a semantically meaningful transition in the skill execution.

task execution through a skill (*e.g.,* a pick and place task would be decomposed as pick $\rightarrow$ place).

Data. We assume access to two datasets:

- $\mathcal{D}_{robot} = \{(o_t, \text{aux}_t, a_t, e_t, g)\}$: Robot teleoperation demonstrations with workspace cameras o_t, auxiliary robot observations aux_t (*e.g.,* wrist cameras, proprioception), actions a_t, end-effector poses e_t and a language goal g, where t denotes the timestep of the observation. We also identify the timesteps where a sub-goal in a demonstration is achieved $s \in \mathcal{S}$ through heuristics (*e.g.,* gripper state changes).
- $\mathcal{D}_{human} = \{(o_t, e_t, g)\}$: Human demonstrations with camera observations o_t and hand poses e_t, derived using an off-the-shelf hand pose estimator [20, 30]. Sub-goal timesteps are manually annotated. Importantly, we note that $\mathcal{D}_{human}$ lacks action labels a_t.

We consider $\mathcal{D}_{robot}$ as **in-distribution** (ID) - demonstrations on objects/tasks for training the low-level policy. $\mathcal{D}_{human}$ contains **out-of-distribution** (OOD) demonstrations of the same skill applied to novel objects, compositions of objects, and contexts where robot data is unavailable.

Assumptions.

- ≥ 1 calibrated RGB-D cameras viewing the workspace for 3D goal prediction.
- The human hand remains visible during the demonstration.
- Human demonstrations use similar grasp types as the robot gripper to minimize the embodiment gap.

Problem Formulation. We learn a hierarchical visuomotor policy $\pi = [\pi_{hi}, \pi_{lo}]$:

- High-level policy $\pi_{hi} : \{o_t, e_t\} \rightarrow e_s$ predicts the end-effector pose e_s at the next sub-goal timestep s, trained on $\mathcal{D}_{robot} \cup \mathcal{D}_{human}$.
- Low-level policy $\pi_{lo} : (o_t, \text{aux}_t, e_s) \rightarrow a_t$ predicts actions conditioned on the goal e_s, trained with imitation learning on $\mathcal{D}_{robot}$.

This decomposition enables π_{lo} to learn precise robot control while π_{hi} takes advantage of human demonstrations for OOD generalization. Note that π_{hi} requires only workspace cameras o_t present in both datasets, allowing π_{lo} to exploit auxiliary sensors aux_t unavailable in human data.

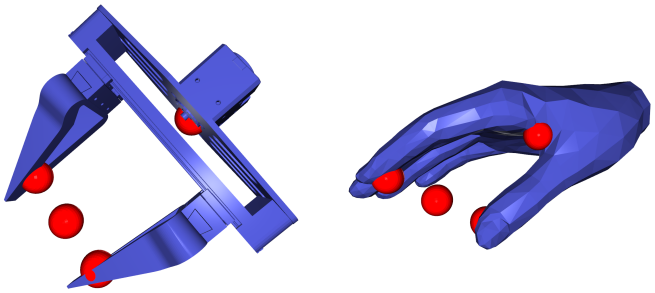

Fig. 4: End-effector representations for robot and human demonstrations. Instead of explicitly retargeting the gripper pose to a translation and SO(3) rotation, we represent gripper poses as sets of 3D keypoints (red spheres): gripper base, fingertip locations, and grasping center. **Left**: Parallel-jaw gripper. **Right**: MANO hand.

IV. GHOST

A. Data Collection and Processing

Robot Data. We collect demonstrations through teleop across N different task variants, each demonstrating the same skill (*e.g.,* pick-and-place), to provide in-distribution coverage. We extract sub-goals $s \in \mathcal{S}$ automatically by identifying the timesteps where the gripper state changes from open $\rightarrow$ close and vice versa.

Human Data. We collect human demonstrations of the *same skill* in a novel setting, using either new objects or novel applications of the skill. We track 3D hand poses H_t using off-the-shelf hand pose estimators [20, 30]. We resolve the weak-perspective scale ambiguity of the detection by segmenting the hand using Grounded-SAM [22] and scaling with the observed depth of the hand. We manually annotate sub-goals at the timesteps where a sub-goal terminates. Figure 3 shows examples of our sub-goal decomposition for both robot and human demonstrations.

Representing the end-effector pose: Following prior work, we represent the end-effector pose as a set of K 3D points instead of a position and SO(3) rotation [26]. For robot data, we sample a point cloud P_R from the gripper mesh at each time-step and represent the end-effector pose e_t as a set of 4 3D points from P_R, located at the base of the gripper, the tips of the two gripper fingers and the grasping center. Similarly, for human data, we extract hand point clouds P_H from the MANO [23] mesh of the hand pose H_t and select 4 3D points - corresponding to points on the palm, tips of the thumb and index finger, and the grasping center. Figure 4 illustrates the keypoint locations for both gripper types. Thus, we extract a unified 3D end-effector representation e_t for human and robot data.

Task	Data Source	# Demos	Generalization Type
Pick-and-Place			
`plate-on-table`	Robot	20	—
`plate-in-bin`	Robot	20	—
`mug-in-bin`	Robot	20	—
`mug-on-table`	Human	20	Compositional (novel object-context combination)
`octopus-on-table`	Human	20	Novel object category
Cloth Folding			
`fold-onesie`	Robot	33	—
`fold-shirt`	Robot	50	—
`fold-onesie-ood`	Human	17	Novel object instance
`fold-onesie-reverse`	Human	25	Novel object instance + Skill composition
`fold-towel`	Human	50	Novel object category + Skill composition
LIBERO-Object			
`LIBERO-Object-9`	Robot	450	—
`LIBERO-Object-heldOut`	Robot	50	Novel object category

TABLE I: Task overview. Blue: ID tasks. Orange: OOD tasks requiring generalization.

B. Policy Learning

1) High-Level Goal Prediction Policy: The high-level policy π_{hi} predicts the end-effector pose $e_s = \{\mathbf{p}_1, \mathbf{p}_2, \mathbf{p}_3, \mathbf{p}_4\}$ at sub-goal time step s, where the 4 points represent 3D positions disambiguating the end-effector [26]. We parameterize π_{hi} as an encoder-only transformer operating on RGB-D observations from C cameras. The architecture is shown in Figure 2.

Each RGB-D image is processed independently with a frozen DINOv2 [16] encoder, producing $256 \times C$ patch tokens. Each patch token is augmented with a 3D coordinate $[x, y, z]$ obtained by transforming the patch-center pixel using the camera parameters and depth. Additional tokens encode the current gripper pose (computed from e_t by extracting the $[x, y, z]$ of the grasping center, the 6D rotation [35] and the width of the gripper) and the task language embedding (with SigLIP [32]). Finally, we also include a learnable token specifying the embodiment that performs the demonstration. All tokens are projected to the transformer's feature dimension through separate MLPs.

To handle the multimodality inherent in the demonstration data [25], we model the goal distributions as a dense per-patch Gaussian Mixture Model (GMM). Each patch token predicts: (1) a mixing weight w_i, and (2) four 3D residual vectors $\{\Delta p_1, \Delta p_2, \Delta p_3, \Delta p_4\}$ relative to the patch center. The global goal distribution is a mixture of $256 \times C$ Gaussians (one per patch) with fixed variance σ^2.

Training. We train π_{hi} by minimizing the negative log-likelihood on sub-goal transitions from $\mathcal{D}' = \mathcal{D}_{robot} \cup \mathcal{D}_{human}$:

$$\mathcal{L}_{hi} = \mathbb{E}_{(o_t, e_s) \sim \mathcal{D}'} \left[-\log \left(\sum_{i=1}^{256 \times C} w_i \mathcal{N}(e_s \mid \mu_i, \sigma^2) \right) \right]$$

where μ_i is the predicted goal (patch center + residuals) for patch i.

Inference. We sample a patch i from the categorical distribution on mixing weights $\{w_i\}$, then use the corresponding mean prediction μ_i as the residual vector to be added to the patch center to obtain the predicted goal.

2) Low-Level Goal Conditioned Policy: We instantiate π_{lo} as a Diffusion Policy [5], which generates chunks of action $\{a_t, ..., a_{t+H}\}$ by denoising conditioned on observations.

Goal Conditioning. The 3D goal e_s is projected onto the image plane of each camera using the camera parameters, yielding sparse 2D coordinates $\mathbf{p}_i \in \mathbb{R}^2$ for $i \in \{1, 2, 3\}$. These are converted to dense *end-effector heatmaps* H where each channel c encodes the pixel distance field from keypoint $\mathbf{p}_c$:

$$H_{xy}^{(c)} = \sqrt{\frac{\|\mathbf{x}_{xy} - \mathbf{p}_c\|_2}{d_{\max}}}$$

where $\mathbf{x}_{xy}$ denotes the pixel coordinates, $d_{\max} = \sqrt{h^2 + w^2}$ normalizes by the image diagonal, and we use the square root for steeper gradients near the targets. The resulting 3-channel heatmap per camera is provided as additional input to the policy via separate image encoders, preserving spatial structure while densifying the sparse goal representation. This process is illustrated in Figure 5.

Training. We use the standard diffusion loss, training on $\mathcal{D}_{robot}$:

$$\mathcal{L}_{lo} = \mathbb{E}_{\mathcal{D}_{robot}, \epsilon} \left[\|\epsilon - \epsilon_\theta(a_t^{\text{noisy}}, o_t, \texttt{aux}_t, \text{heatmap}(e_s))\|^2 \right]$$

We train π_{hi} and π_{lo} independently, using ground-truth goals from demonstrations e_s when training π_{lo}.

Inference. At deployment, we execute π_{hi} and π_{lo} sequentially, where π_{hi} predicts a subgoal conditioned on the current observation and π_{lo} predicts one action chunk conditioned on the goal prediction.

V. Experiments and Results

We evaluate whether human video enables skill generalization in hierarchical imitation learning. Our experiments test three hypotheses: (1) hierarchical policies improve in-distribution performance even without human data, (2) human demonstrations enable compositional generalization (novel

Method	Success Rate (%)		
	`plate-on-table`	`mug-on-table`	`octopus-on-table`
DP	**60**	**5**	0
Hierarchical DP w/out Human	**70**	0	0
GHOST (Ours)	**60**	**32.5**	**20**

TABLE II: Pick-and-place results. Best: **shaded**, second-best: **bold**. Blue: ID. Orange: OOD.

Task	Method	Success Rate (%)				
		1-step	2-step	3-step	4-step	Final
`fold-onesie`	DP	**100**	70	30	0	0
	Hierarchical DP w/out Human	**100**	**100**	**100**	**100**	**80**
	GHOST (Ours)	**100**	**100**	**90**	**90**	**70**
`fold-onesie-ood` (Novel Instance)	DP	40	40	10	0	0
	Hierarchical DP w/out Human	**100**	**100**	60	50	40
	GHOST (Ours)	**100**	**100**	**90**	**70**	**50**
`fold-onesie-reverse` (Novel Composition)	GHOST (Ours)	100	100	70	20	0
`fold-towel` (Novel Category)	GHOST (Ours)	60	—	—	—	5

TABLE III: Cloth folding results. k-step = first k folds completed. Final = all folds completed (5-step for the **`fold-onesie`** tasks, 2-step for **`fold-towel`**). Best: **shaded**, second-best: **bold**. Blue: ID. Orange: OOD.

Fig. 5: End-effector heatmap generation for goal conditioning. Predicted 3D keypoints are projected to 2D coordinates, and then converted to dense distance field heatmaps that encode spatial proximity to each keypoint.

Method	Success Rate (%)	
	`LIBERO-Object-9`	`LIBERO-Object-heldOut`
DP	**88**	0
Hierarchical DP	**90**	**24**

TABLE IV: LIBERO-Object results. Best: **shaded**, second-best: **bold**. Blue: ID. Orange: OOD.

object-context combinations and novel skill combinations), and (3) learned skills generalize to novel object instances and categories.

A. Tasks and Evaluation Protocol

Table I summarizes our evaluation tasks. We collect robot demonstrations on in-distribution (ID) tasks to train both π_{hi} and π_{lo}, and human demonstrations on out-of-distribution (OOD) variations to train only π_{hi}. For LIBERO-Object [13], we use robot demonstrations for both the ID and OOD tasks.

- **Pick-and-Place:** Grasp a specified object (plate, mug, octopus) and place on a target (table, bin). Robot demos cover plate and mug on the two targets; human demos provide OOD object-target combinations. The task is considered a success if the object is placed stably on the target, objects placed successfully but knocked off during reset are given a score of 0.5.
- **Cloth Folding:** Long-horizon deformable manipulation requiring multiple sequential folds. Robot demos on folding onesies and shirts; human demos on novel instances (**`fold-onesie-ood`**), composing the skill in novel ways (**`fold-onesie-reverse`**), and novel categories entirely (**`fold-towel`**). k-step success indicates first k folds completed successfully
- **LIBERO-Object:** Standardized pick-and-place benchmark [13] with 10 objects. We train π_{lo} on 9 objects, π_{hi} on all 10, and evaluate transfer to the held-out object. Success is reported by the simulation environment.

B. Baselines

We evaluate GHOST against two baselines: (1) Diffusion Policy (DP) [5], a flat policy trained on robot data, and (2) Hierarchical DP, our hierarchical framework trained only on robot data without human demonstrations. All policies are trained as multi-task policies with language conditioning. For our hierarchical framework, language conditioning is only provided for the high-level policy. For the Diffusion Policy baseline, we add language conditioning by appending the text embedding to the global conditioning vector.

C. Implementation

We train all DP variants for 300k steps with color jitter and random crop augmentations. The high-level policy is trained for 100 epochs with additional augmentations: adding

Qualitative Results

Fig. 6: Qualitative Results

noise to the gripper pose, blur, grayscale, and token dropout augmentations. At inference, we synchronously execute π_{hi} and π_{lo} - the low-level predicts action chunks that fully execute before the next high-level inference step.

D. Results

Sriram: [Work in progress...]

Pick-and-Place (Table II): GHOST achieves comparable ID performance (60%) while enabling substantial OOD generalization (32.5% on mug-on-table, 20% on octopus-on-table). Hierarchical DP w/out Human fails completely on OOD tasks despite slightly better ID performance (70%), confirming hierarchy alone is insufficient and the human demos are critical for generalizing goal prediction.

Cloth Folding (Table III): On ID fold-onesie, Hierarchical DP w/out Human achieves the best final success (80%), with GHOST second (70%). However, GHOST substantially outperforms on OOD tasks: 50% final success on fold-onesie-ood vs 40% for Hierarchical DP and 0% for flat DP. Flat DP collapses on long-horizon tasks without hierarchy. GHOST demonstrates compositional transfer on fold-onesie-reverse (novel skill sequencing) and fold-towel (novel object category), though performance degrades significantly as the horizon of the task increases (5% final success).

LIBERO-Object (Table IV): Hierarchical DP achieves 24% OOD success vs 0% for flat DP.

VI. Limitations and Future Work

Although GHOST shows promising results in generalizing learned skills across multiple axes in a sample-efficient setting, there are numerous limitations. Like other methods that use hand pose estimators on human demonstrations [12, 21], our approach is limited by the quality of the estimated hand pose, and suffers in situations where hand pose estimation is weak, such as occlusions.

We extract sub-goals through heuristics for robot data (gripper open/close), which may not be optimal or even valid for tasks requiring continuous manipulation skills (pouring, stirring), and through annotation on human demonstrations, which imposes an additional cost on the collection of demonstrations. Related work on automatic sub-goal discovery [19, 33] is an interesting direction for future research on mining sub-goals from collected human and robot demonstrations.

While GHOST is able to generalize to novel object instances and compose learned skills in novel sequences, we observe a steep drop in performance as the horizon of the task lengthens, which we observe to be caused by the visual domain gap for the high-level policy. Existing work using human and robot data seeks to learn embodiment-invariant visual representations [25, 21] to minimize this gap and represents an important research direction for future work.

Lastly, the promise of this hierarchical policy factorization lies in its ability to scale high-level and low-level policies independently on large human and robot datasets. This allows

learning both generalizable goal prediction and goal-following behavior, and is an exciting path forward towards learning generalist robot policies.

VII. Conclusion

We present GHOST, a hierarchical imitation learning framework that decouples embodiment-agnostic goal prediction from embodiment-specific action execution. This factorization enables two key benefits: improved in-distribution performance through explicit sub-goal modeling, and out-of-distribution generalization via human demonstrations without requiring action labels from human data.

We introduce a 3D goal prediction architecture that processes multi-view RGB-D observations and predicting dense per-patch Gaussian mixture models over end-effector poses. We represent goals as unified 3D keypoint sets across embodiments and condition the low-level policy through end-effector heatmaps. This representation bridges the gap between human and robot demonstrations without requiring explicit action retargeting.

By training π_{hi} on heterogeneous human and robot data and π_{lo} purely on robot demonstrations, we achieve compositional generalization across novel object-context combinations (32.5% success versus 0% for baselines on out-of-distribution pick-and-place) and meaningful progress on novel skill compositions (60% task progress on novel cloth folding instances).

Acknowledgments

We would like to thank Alexis Hao, Mino Nakura, Kallol Saha and Pratik Bhowal for helpful discussions and assistance with data collection. We also thank the other members of the R-PAD lab for their feedback.

References

[1] Jose Barreiros, Andrew Beaulieu, Aditya Bhat, Rick Cory, Eric Cousineau, Hongkai Dai, Ching-Hsin Fang, Kunimatsu Hashimoto, Muhammad Zubair Irshad, Masha Itkina, et al. A careful examination of large behavior models for multitask dexterous manipulation. *arXiv preprint arXiv:2507.05331*, 2025.

[2] Homanga Bharadhwaj, Roozbeh Mottaghi, Abhinav Gupta, and Shubham Tulsiani. Track2act: Predicting point tracks from internet videos enables diverse zero-shot robot manipulation. *CoRR*, 2024.

[3] Kevin Black, Noah Brown, Danny Driess, Adnan Esmail, Michael Equi, Chelsea Finn, Niccolo Fusai, Lachy Groom, Karol Hausman, Brian Ichter, et al. π_0: A vision-language-action flow model for general robot control. *arXiv preprint arXiv:2410.24164*, 2024.

[4] Anthony Brohan, Noah Brown, Justice Carbajal, Yevgen Chebotar, Joseph Dabis, Chelsea Finn, Keerthana Gopalakrishnan, Karol Hausman, Alex Herzog, Jasmine Hsu, et al. Rt-1: Robotics transformer for real-world control at scale. *arXiv preprint arXiv:2212.06817*, 2022.

[5] Cheng Chi, Zhenjia Xu, Siyuan Feng, Eric Cousineau, Yilun Du, Benjamin Burchfiel, Russ Tedrake, and Shuran Song. Diffusion policy: Visuomotor policy learning via action diffusion. *The International Journal of Robotics Research*, 44(10-11):1684–1704, 2025.

[6] Jeremy A Collins, Loránd Cheng, Kunal Aneja, Albert Wilcox, Benjamin Joffe, and Animesh Garg. Amplify: Actionless motion priors for robot learning from videos. *arXiv preprint arXiv:2506.14198*, 2025.

[7] Raghav Goyal, Samira Ebrahimi Kahou, Vincent Michalski, Joanna Materzynska, Susanne Westphal, Heuna Kim, Valentin Haenel, Ingo Fruend, Peter Yianilos, Moritz Mueller-Freitag, et al. The" something something" video database for learning and evaluating visual common sense. In *Proceedings of the IEEE international conference on computer vision*, pages 5842–5850, 2017.

[8] Kristen Grauman, Andrew Westbury, Eugene Byrne, Zachary Chavis, Antonino Furnari, Rohit Girdhar, Jackson Hamburger, Hao Jiang, Miao Liu, Xingyu Liu, et al. Ego4d: Around the world in 3,000 hours of egocentric video. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 18995–19012, 2022.

[9] Siddhant Haldar and Lerrel Pinto. Point policy: Unifying observations and actions with key points for robot manipulation. *arXiv preprint arXiv:2502.20391*, 2025.

[10] Simar Kareer, Dhruv Patel, Ryan Punamiya, Pranay Mathur, Shuo Cheng, Chen Wang, Judy Hoffman, and Danfei Xu. Egomimic: Scaling imitation learning via egocentric video. In *2025 IEEE International Conference on Robotics and Automation (ICRA)*, pages 13226–13233. IEEE, 2025.

[11] Alexander Khazatsky, Karl Pertsch, Suraj Nair, Ashwin Balakrishna, Sudeep Dasari, Siddharth Karamcheti, Soroush Nasiriany, Mohan Kumar Srirama, Lawrence Yunliang Chen, Kirsty Ellis, et al. Droid: A large-scale in-the-wild robot manipulation dataset. *arXiv preprint arXiv:2403.12945*, 2024.

[12] Marion Lepert, Jiaying Fang, and Jeannette Bohg. Phantom: Training robots without robots using only human videos. *arXiv preprint arXiv:2503.00779*, 2025.

[13] Bo Liu, Yifeng Zhu, Chongkai Gao, Yihao Feng, Qiang Liu, Yuke Zhu, and Peter Stone. Libero: Benchmarking knowledge transfer for lifelong robot learning. *Advances in Neural Information Processing Systems*, 36:44776–44791, 2023.

[14] Corey Lynch, Mohi Khansari, Ted Xiao, Vikash Kumar, Jonathan Tompson, Sergey Levine, and Pierre Sermanet. Learning latent plans from play. In *Conference on robot learning*, pages 1113–1132. Pmlr, 2020.

[15] Ajay Mandlekar, Danfei Xu, Roberto Martín-Martín, Silvio Savarese, and Li Fei-Fei. Learning to generalize across long-horizon tasks from human demonstrations. *arXiv preprint arXiv:2003.06085*, 2020.

[16] Maxime Oquab, Timothée Darcet, Théo Moutakanni, Huy Vo, Marc Szafraniec, Vasil Khalidov, Pierre Fernandez, Daniel Haziza, Francisco Massa, Alaaeldin El-Nouby, et al. Dinov2: Learning robust visual features

without supervision. *arXiv preprint arXiv:2304.07193*, 2023.

[17] Abby O'Neill, Abdul Rehman, Abhiram Maddukuri, Abhishek Gupta, Abhishek Padalkar, Abraham Lee, Acorn Pooley, Agrim Gupta, Ajay Mandlekar, Ajinkya Jain, et al. Open x-embodiment: Robotic learning datasets and rt-x models: Open x-embodiment collaboration 0. In *2024 IEEE International Conference on Robotics and Automation (ICRA)*, pages 6892–6903. IEEE, 2024.

[18] Georgios Pavlakos, Dandan Shan, Ilija Radosavovic, Angjoo Kanazawa, David Fouhey, and Jitendra Malik. Reconstructing hands in 3d with transformers. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 9826–9836, 2024.

[19] Karl Pertsch, Oleh Rybkin, Frederik Ebert, Shenghao Zhou, Dinesh Jayaraman, Chelsea Finn, and Sergey Levine. Long-horizon visual planning with goal-conditioned hierarchical predictors. *Advances in Neural Information Processing Systems*, 33:17321–17333, 2020.

[20] Rolandos Alexandros Potamias, Jinglei Zhang, Jiankang Deng, and Stefanos Zafeiriou. Wilor: End-to-end 3d hand localization and reconstruction in-the-wild. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pages 12242–12254, 2025.

[21] Juntao Ren, Priya Sundaresan, Dorsa Sadigh, Sanjiban Choudhury, and Jeannette Bohg. Motion tracks: A unified representation for human-robot transfer in few-shot imitation learning. *arXiv preprint arXiv:2501.06994*, 2025.

[22] Tianhe Ren, Shilong Liu, Ailing Zeng, Jing Lin, Kunchang Li, He Cao, Jiayu Chen, Xinyu Huang, Yukang Chen, Feng Yan, et al. Grounded sam: Assembling open-world models for diverse visual tasks. *arXiv preprint arXiv:2401.14159*, 2024.

[23] Javier Romero, Dimitrios Tzionas, and Michael J. Black. Embodied hands: Modeling and capturing hands and bodies together. *ACM Transactions on Graphics, (Proc. SIGGRAPH Asia)*, 36(6), November 2017.

[24] Aaron Van Den Oord, Oriol Vinyals, et al. Neural discrete representation learning. *Advances in neural information processing systems*, 30, 2017.

[25] Chen Wang, Linxi Fan, Jiankai Sun, Ruohan Zhang, Li Fei-Fei, Danfei Xu, Yuke Zhu, and Anima Anandkumar. Mimicplay: Long-horizon imitation learning by watching human play. *arXiv preprint arXiv:2302.12422*, 2023.

[26] Yufei Wang, Ziyu Wang, Mino Nakura, Pratik Bhowal, Chia-Liang Kuo, Yi-Ting Chen, Zackory Erickson, and David Held. Articubot: Learning universal articulated object manipulation policy via large scale simulation. *arXiv preprint arXiv:2503.03045*, 2025.

[27] Chuan Wen, Xingyu Lin, John So, Kai Chen, Qi Dou, Yang Gao, and Pieter Abbeel. Any-point trajectory modeling for policy learning. *arXiv preprint arXiv:2401.00025*, 2023.

[28] Danfei Xu, Suraj Nair, Yuke Zhu, Julian Gao, Animesh Garg, Li Fei-Fei, and Silvio Savarese. Neural task programming: Learning to generalize across hierarchical tasks. In *2018 IEEE international conference on robotics and automation (ICRA)*, pages 3795–3802. IEEE, 2018.

[29] Seonghyeon Ye, Joel Jang, Byeongguk Jeon, Sejune Joo, Jianwei Yang, Baolin Peng, Ajay Mandlekar, Reuben Tan, Yu-Wei Chao, Bill Yuchen Lin, et al. Latent action pretraining from videos. *arXiv preprint arXiv:2410.11758*, 2024.

[30] Yufei Ye, Yao Feng, Omid Taheri, Haiwen Feng, Shubham Tulsiani, and Michael J Black. Predicting 4d hand trajectory from monocular videos. *arXiv preprint arXiv:2501.08329*, 2025.

[31] Yanjie Ze, Gu Zhang, Kangning Zhang, Chenyuan Hu, Muhan Wang, and Huazhe Xu. 3d diffusion policy: Generalizable visuomotor policy learning via simple 3d representations. *arXiv preprint arXiv:2403.03954*, 2024.

[32] Xiaohua Zhai, Basil Mustafa, Alexander Kolesnikov, and Lucas Beyer. Sigmoid loss for language image pre-training. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 11975–11986, 2023.

[33] Zichen Zhang, Yunshuang Li, Osbert Bastani, Abhishek Gupta, Dinesh Jayaraman, Yecheng Jason Ma, and Luca Weihs. Universal visual decomposer: Long-horizon manipulation made easy. In *2024 IEEE International Conference on Robotics and Automation (ICRA)*, pages 6973–6980. IEEE, 2024.

[34] Tony Z Zhao, Vikash Kumar, Sergey Levine, and Chelsea Finn. Learning fine-grained bimanual manipulation with low-cost hardware. *arXiv preprint arXiv:2304.13705*, 2023.

[35] Yi Zhou, Connelly Barnes, Jingwan Lu, Jimei Yang, and Hao Li. On the continuity of rotation representations in neural networks. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 5745–5753, 2019.